

Comprehensive Report: Fundamentals of Generative AI and Large Language Models

Name: Harish Kumar P

Register no:212225230095

Slot:T2-H16

Subject:19TD608- PROMPT ENGINEERING

1.1 Subject: Generative AI, LLMs, and 2024-2025 AI Tool Landscape

1.2 Target Audience: Students and Technical Professionals

1.3 Table of Contents

1. Introduction

2. Introduction to Artificial Intelligence and Machine Learning

3. What is Generative AI?

4. Generative Models and Their Types

5.2024 AI Tools

6. Introduction to Large Language Models (LLMs)

7. Architecture of Large Language Models

8. Training Process and Data Requirements

G.Use Cases and Applications

10.Limitations and Ethical Considerations

1. Executive Summary

This report provides a foundational overview of Generative AI, tracing its evolution from early machine learning to the sophisticated Large Language Models (LLMs) of today. It explores the technical mechanisms behind model training, highlights key tools from the 2024 landscape, and outlines the architectural shifts—most notably the Transformer—that made modern AI possible.

2. Introduction to Generative AI

Artificial Intelligence (AI) has evolved from simple rule-based systems to advanced models capable of generating human-like text, images, audio, and videos. Generative AI and Large Language Models are at the core of this transformation. This report aims to explain these technologies in a clear and structured manner for students and beginners. Generative AI is a branch of Artificial Intelligence focused on creating new content. While traditional AI (Discriminative AI) is designed to recognize patterns and classify data (e.g., "Is this a picture of a cat?"), Generative AI learns the underlying distribution of data to generate new, similar examples (e.g., "Create a new picture of a cat")

2.1 Foundational Concepts Generative Model:

A statistical model that acts as a "creator." It models how the data was generated in order to categorize a signal or create new data points.

Probability Distribution: The mathematical framework the model uses to understand the likelihood of a certain pixel or word following another.

Introduction to Artificial Intelligence and Machine Learning

Artificial Intelligence (AI) refers to the simulation of human intelligence in machines that are programmed to think, learn, and make decisions.

Machine Learning (ML) is a subset of AI that allows systems to learn patterns from data without being explicitly programmed. Types of Machine Learning:

Supervised Learning - Learning from labeled data

Unsupervised Learning - Finding patterns in unlabeled data

Reinforcement Learning - Learning through rewards and penalties

3. What is Generative AI?

Generative AI is a class of artificial intelligence models that can generate new content such as text, images, music, code, and videos. Unlike traditional AI systems that only analyze or classify data, Generative AI creates original outputs based on learned patterns.

Key Characteristics:

- **Learns data distributions**
- **Produces new and realistic content**
- **Uses probability and deep learning**

Examples include ChatGPT (text), DALL·E (images), and MusicLM (audio).

3.2Types of Generative Models

2.2 Types of Generative Models		
Model Type	Mechanism	Primary Use Case
GANs (Generative Adversarial Networks)	Two networks (Generator and Discriminator) compete. One creates, the other critiques.	High-quality image generation, deepfakes.
VAEs (Variational Autoencoders)	Encodes input into a compressed space and decodes it back to create variations.	Image denoising, molecular discovery.
Diffusion Models	Gradually adds noise to data and then learns to reverse the process to "recover" an image.	Midjourney, DALL-E 3, Stable Diffusion.

3.3Generative Models and Their Types:

A Generative Model learns the underlying distribution of data and generates new samples similar to the training data.

4. Major Types of Generative Models:

4.1 Generative Adversarial Networks (GANs)

- Consist of two networks: Generator and Discriminator
- Generator creates fake data
- Discriminator evaluates real vs fake data
- Used in image generation and deepfakes

4.2 Variational Autoencoders (VAEs)

- Encode data into a latent space
- Decode it back to generate new data
- Useful for image reconstruction and anomaly detection

4.3 Diffusion Models

- Generate data by gradually removing noise
- Produce high-quality images
- Used in Stable Diffusion and DALL·E 2

5.The 2024 AI Toolscape

The year 2024 marked a peak in multimodal capabilities, where tools began handling text, image, and video simultaneously with high fidelity.

Text G Reasoning: GPT-4o (OpenAI), Claude 3.5 Sonnet (Anthropic), Gemini 1.5 Pro (Google).

Image Generation: Midjourney v6, DALL-E 3, Adobe Firefly.

Video Generation: Sora (OpenAI), Runway Gen-3, Luma Dream Machine.

Coding: GitHub Copilot, Cursor, Replit Ghostwriter.

5.1 2024 AI Tools Some popular AI tools used in 2024 include:

Tool Name	Purpose
ChatGPT	Text generation, tutoring
Google Gemini	Multimodal AI
Claude	Long-document reasoning
Midjourney	Image generation
Stable Diffusion	Open-source image generation
GitHub Copilot	Code assistance
Canva AI	Design and content creation

6. Large Language Models (LLMs): How They are Built

An LLM is a massive neural network trained to predict the "next token" (word or sub-word) in a sequence.

6.1 The Building Blocks:

Transformer Architecture Self-Attention Mechanism: Allows the model to weigh the importance of different words in a sentence simultaneously.

Positional Encoding: Since Transformers process all words at once, they use mathematical codes to remember the order of words.

6.2 The Training Process:

Pre-training: The model "reads" trillions of words from the internet to learn grammar, facts, and reasoning.


Supervised Fine-Tuning (SFT): The model is trained on specific prompt-response pairs to learn how to follow instructions.

RLHF (Reinforcement Learning from Human Feedback): Humans rank model responses to align the AI with human values (helpfulness and safety).

LLMs can:

- Answer questions
- Summarize text
- Translate languages
- Generate code Examples: GPT-3, GPT-4, BERT, PaLM

7. Architecture of Large Language Models



7.1 Transformer Architecture

The transformer is the backbone of modern LLMs.

Key components:

Tokenization Embedding Layer Self-Attention Mechanism Feed-Forward Neural
Networks Output Layer

7.2 Popular LLM Architectures

- GPT (Generative Pre-trained Transformer) - Decoder-only model
- BERT (Bidirectional Encoder Representations from Transformers) - Encoder-only model

8. Training Process and Data Requirements

Training Steps:

- Data Collection (books, websites, articles)
- Data Cleaning and Preprocessing
- Pre-training on large datasets
- Fine-tuning for specific tasks
- Evaluation and deployment

Data Requirements:

- Large-scale text data
- High-quality and diverse sources
- Powerful hardware (GPUs/TPUs)

G. Use Cases and Applications

Major Applications:

- Chatbots and virtual assistants
- Content writing and summarization
- Code generation and debugging
- Healthcare and diagnostics
- Education and personalized learning
- Business analytics and automation

10. Technical Implementation:

Next-Token Prediction Logic

LLMs function essentially as advanced probability machines. Below is a conceptual implementation of how "Self-Attention" weights might be calculated.

```
import numpy as np
```

```
def calculate_attention(query, key, value):
```

Step 1: Compute attention scores (Dot Product)

```
scores = np.dot(query, key.T)
```

Step 2: Normalize with Softmax (sum of weights = 1)

```
weights = np.exp(scores) / np.sum(np.exp(scores))
```

Step 3: Weighted sum of values

```
output = np.dot(weights, value)
```

```
return output, weights
```

Example: Finding context for the word "Bank" (River bank vs Money bank)

```
q = np.array([1, 0, 0]) # "Query" for river context
```

```
k = np.array([1, 0, 0]) # "Key" for the word "Water"
```

```
v = np.array([10, 20, 30]) # "Value" containing semantic data
```

```
context_result, attention_weights = calculate_attention(q, k, v)
```

```
print(f'Context Vector: {context_result}')
```

11. Limitations and Ethical Considerations

Limitations:

1. High computational cost
2. Data bias
3. Hallucinations (incorrect outputs)
4. Limited real-world understanding

Ethical Issues:

1. Privacy concerns
2. Misinformation

3. Job displacement


4. Copyright and data ownership

12. Future Trends in Generative AI and LLMs

- Multimodal AI (text + image + audio)
- Smaller and efficient models
- Personalized AI assistants
- Better alignment and safety
- Increased use in education and healthcare

13. Evolution of AI - Timeline Chart

The following chart traces the journey from early logic to modern generative power:



1950: Alan Turing proposes the "Turing Test."

1956: Dartmouth Conference; the term "Artificial Intelligence" is coined.

1966: ELIZA, the first chatbot, is created at MIT.

1997: IBM's Deep Blue defeats world chess champion Garry Kasparov.

2012: AlexNet triggers the "Deep Learning" revolution.

2017: Google publishes "Attention is All You Need" (The birth of Transformers).

2022: OpenAI launches ChatGPT, bringing Generative AI to the mainstream.

2024-2026: Rise of Multimodal "Omni" models and Autonomous AI Agents.

14. Conclusion

Generative AI and Large Language Models represent a revolutionary advancement in artificial intelligence. By understanding their foundations, architectures, and applications, students can better prepare for future careers in AI-driven industries. While challenges and ethical concerns exist, responsible development and usage can unlock enormous potential. The rapid evolution from the basic logic of the 1950s to the generative capabilities of 2026 highlights a shift toward AI that understands context and creative intent. As LLMs become more multimodal and efficient, the focus is now moving toward Agentic AI—systems that can not only generate text but execute complex, multi-step tasks autonomously.

8. References

- OpenAI Documentation
- Google AI Research Blogs
- Deep Learning by Ian Goodfellow
- Research Papers on Transformers and LLMs
- Online AI educational resources
- Vaswani, A., et al. (2017). Attention Is All You Need.
- OpenAI (2024). GPT-4o Technical Report.
- Goodfellow, I., et al. (2014). Generative Adversarial Networks.